

Lane Line Detection Using OpenCV

This document describes the Python script used to detect lane lines in images and videos using OpenCV, an open-source computer vision library.

First, let's import the following python libraries.

```
1 import cv2
2 import numpy as np
3 import matplotlib.pyplot as plt
```

Next, let's import an image with lane lines in order to develop our script.

```
1 #Load and display lane lines image
2 image = cv2.imread('test_image.jpg')
3 cv2.imshow('result', image)
4 cv2.waitKey(0)
```

-1



As we will be using a method called **Canny Edge detection** which is not concerned about colours, we can first **convert the image to grayscale**.

```
1 #make a copy of image and convert it to grayscale
2 lane_image = np.copy(image)
3 gray = cv2.cvtColor(lane_image, cv2.COLOR_RGB2GRAY)
4 cv2.imshow('result', gray)
5 cv2.waitKey(0)
```



-1




Since edge detection is susceptible to noise in the image, the next step is to **remove the noise in the image with a 5x5 Gaussian filter**. Gaussian blur works by smoothing / averaging out the intensity of a pixel by its surrounding pixels' intensity.

```
1 #filter out image noise (average out the intensity of a pixel by its surrounding pixels' intensity)
2 blur = cv2.GaussianBlur(gray, (5,5), 0)
3 cv2.imshow('result', blur)
4 cv2.waitKey(0)
```

Example of how smoothing affects edge detection






Now, we can apply the **cv2.Canny()** method to the image to determine areas with substantial change in pixel intensity (e.g. black to white). These are edges with steep gradient of intensity.

```
1 #determine the areas with substantial change in intensity (steep gradient e.g. black to white)
2 canny = cv2.Canny(blur, 50, 150)
3 cv2.imshow('result', canny)
4 cv2.waitKey(0)
```

-1




We can combine the above functions into 1 and call it canny().


```
1 #summarising the above codes into a function and using it
2 def canny(image):
3     gray = cv2.cvtColor(lane_image, cv2.COLOR_RGB2GRAY)
4     blur = cv2.GaussianBlur(gray, (5,5), 0)
5     canny = cv2.Canny(blur, 50, 150)
6     return canny
```

Next, we want to **filter out the region of interest** which contains the lane lines. To do this, we can use the function below. What it basically does is to **create an image with a white colour triangular region that surrounds the lane lines**. Everything else outside the region will be black. We then overlay that region over the original image and return areas where both layers are white.

```
def region_of_interest(image):  
    height = image.shape[0]                                #height should be 704 pixels  
    polygons = np.array([                                  #define a triangular area of interest  
        [(200, height), (1100, height), (550,250)]  
    ])   
    mask = np.zeros_like(image)                            #create an array of zeros (all black image) with same shape as image  
    cv2.fillPoly(mask, polygons, 255)                     #fill up the triangular area of interest with white colour  
    masked_image = cv2.bitwise_and(image, mask)            #overlay the area of interest on the image  
    return masked_image
```



After Filtering Out Region of Interest



```

52 image = cv2.imread('test_image.jpg')
53 lane_image = np.copy(image)
54 canny_image = canny(lane_image)
55 cropped_image = region_of_interest(canny_image)

```

With this, we can now use the **cv2.HoughLinesP()** method to detect straight lines. Note that this function outputs the extremes of the detected lines (x0,y0,x1,y1). Refer to code comments below for details of the function parameters.

```

57 #identify lines in the cropped_image down to accuracy of 2 pixels, 1 degree, and 100 intersections in each bin.
58 #each line must minimally be 40 pixels in length and if gap between lines are less than 5 pixels, they will combine into 1.
59 lines = cv2.HoughLinesP(cropped_image, 2, np.pi/180, 100, np.array([]), minLineLength=40, maxLineGap=5)

```


Now that we obtain the coordinates of the lines, we can use the following method to **display the lines**.

```

34 def display_lines(image, lines):
35     line_image = np.zeros_like(image)
36     if lines is not None:
37         for line in lines:
38             x1, y1, x2, y2 = line.reshape(4) #reshape 2D to 1D array
39             cv2.line(line_image, (x1, y1), (x2, y2), (255, 0, 0), 10) #only show the lines (blue colour, thickness = 10) i
40     return line_image
..

60 averaged_lines = average_slope_intercept(lane_image, lines) #average the lines on both sides (display only 1 line each side
61 line_image = display_lines(lane_image, averaged_lines)
62 cv2.imshow('result', line_image)
63 cv2.waitKey(0)

```



Next, we will use the `cv2.addWeighted()` function to overlay the line on top of our original image.

```
1 #display the blue line right on top of the original image
2 combo_image = cv2.addWeighted(lane_image, 0.8, line_image, 1, 1)
3 cv2.imshow('result', combo_image)
4 cv2.waitKey(0)
```

-1



What's remaining is to **do the same thing for a video** instead of a single image. To do that, we just need to use `cv2.VideoCapture()` to read individual frames, and then use back the above codes to overlay the lines on top of the frames.

```
1 #capture the video
2 cap = cv2.VideoCapture('test2.mp4')
3 while (cap.isOpened()):
4     _, frame = cap.read()
5     #reuse the above script (change lane_image to frame)
6     canny_image = canny(frame)
7     cropped_image = region_of_interest(canny_image)
8     lines = cv2.HoughLinesP(cropped_image, 2, np.pi/180, 100, np.array([]), minLineLength=40, maxLineGap=5)
9     averaged_lines = average_slope_intercept(frame, lines)
10    line_image = display_lines(frame, averaged_lines)
11    combo_image = cv2.addWeighted(frame, 0.8, line_image, 1, 1)
12    cv2.imshow('result', combo_image)
13    if cv2.waitKey(1) & 0xFF == ord('q'):
14        break
15 cap.release()
16 cv2.destroyAllWindows()
```

