

## Traffic Signs Recognition using Convolutional Neural Network (CNN)

This document describes the Python script used to generate and train a CNN model that is capable of recognising German road traffic signs with close to 97% accuracy.

First, we import the relevant modules. The most important one to note here is the Keras library, a high-level, deep learning API developed by Google for easy implementation of neural networks.

```
✓ 1s ▶ import numpy as np
import matplotlib.pyplot as plt
import keras
from keras.models import Sequential
from keras.optimizers import Adam
from keras.layers import Dense
from keras.layers import Flatten, Dropout
from keras.utils.np_utils import to_categorical
from keras.layers.convolutional import Conv2D, MaxPooling2D
import random
import pickle
import pandas as pd
import cv2

from keras.callbacks import LearningRateScheduler, ModelCheckpoint
%matplotlib inline

✓ 0s [4] np.random.seed(0)
```

Next, we clone the data available on bitbucket into google drive so that we can load the datasets.

```
✓ 13s [1] !git clone https://bitbucket.org/jadslim/german-traffic-signs

Cloning into 'german-traffic-signs'...
Unpacking objects: 100% (6/6), done.

✓ 0s [2] !ls german-traffic-signs

signnames.csv  test.p  train.p  valid.p
```

We can now load the training, validation, and testing datasets. Because the datasets were saved in pickle format, we need to read the file in binary format ('rb') and use the Pickle module to load the file.

Once we have done that, we can save the features and labels into different variables and check that the shapes of the variables are correct.

We can also read the signnames.csv using Pandas library – read\_csv() method and save the dataframe into a variable called **data**.

```
[5] #load the data
with open('german-traffic-signs/train.p', 'rb') as f:
    train_data = pickle.load(f)
with open('german-traffic-signs/valid.p', 'rb') as f:
    val_data = pickle.load(f)
#load test data
with open('german-traffic-signs/test.p', 'rb') as f:
    test_data = pickle.load(f)

print(type(train_data))
```

```
<class 'dict'>
```

```
✓ [6] # Split out features and labels
Ds
X_train, y_train = train_data['features'], train_data['labels']
X_val, y_val = val_data['features'], val_data['labels']
X_test, y_test = test_data['features'], test_data['labels']
```

```
✓ [7] #4 dimensional (rgb data)
Ds
print(X_train.shape)
print(X_val.shape)
print(X_test.shape)

#check if shape of data is as expected
assert(X_train.shape[0] == y_train.shape[0]), "The number of images is not equal to the number of labels."
assert(X_train.shape[1:] == (32,32,3)), "The dimensions of the images are not 32 x 32 x 3."
assert(X_val.shape[0] == y_val.shape[0]), "The number of images is not equal to the number of labels."
assert(X_val.shape[1:] == (32,32,3)), "The dimensions of the images are not 32 x 32 x 3."
assert(X_test.shape[0] == y_test.shape[0]), "The number of images is not equal to the number of labels."
assert(X_test.shape[1:] == (32,32,3)), "The dimensions of the images are not 32 x 32 x 3."
data = pd.read_csv('german-traffic-signs/signnames.csv')

(34799, 32, 32, 3)
(4410, 32, 32, 3)
(12630, 32, 32, 3)
```

Now, let's print out the dataframe. We can see that the dataframe consists of 2 fields, first being the index of the traffic sign and second being the name of the sign. There are a total of **43 traffic sign categories**.



```
print(data)
```



	ClassId	SignName
0	0	Speed limit (20km/h)
1	1	Speed limit (30km/h)
2	2	Speed limit (50km/h)
3	3	Speed limit (60km/h)
4	4	Speed limit (70km/h)
5	5	Speed limit (80km/h)
6	6	End of speed limit (80km/h)
7	7	Speed limit (100km/h)
8	8	Speed limit (120km/h)
9	9	No passing
10	10	No passing for vechiles over 3.5 metric tons
11	11	Right-of-way at the next intersection
12	12	Priority road
13	13	Yield
14	14	Stop
15	15	No vechiles
16	16	Vechiles over 3.5 metric tons prohibited
17	17	No entry
18	18	General caution
19	19	Dangerous curve to the left
20	20	Dangerous curve to the right
21	21	Double curve
22	22	Bumpy road
23	23	Slippery road
24	24	Road narrows on the right
25	25	Road work
26	26	Traffic signals
27	27	Pedestrians
28	28	Children crossing
29	29	Bicycles crossing
30	30	Beware of ice/snow
31	31	Wild animals crossing
32	32	End of all speed and passing limits
33	33	Turn right ahead
34	34	Turn left ahead
35	35	Ahead only
36	36	Go straight or right
37	37	Go straight or left
38	38	Keep right
39	39	Keep left
40	40	Roundabout mandatory
41	41	End of no passing
42	42	End of no passing by vechiles over 3.5 metric ...

For each of the category, let's randomly display 5 images so that we can better understand what our data contains.

```
✓ [9] num_of_samples=[]  
17s  
  
cols = 5  
num_classes = 43  
  
fig, axs = plt.subplots(nrows=num_classes, ncols=cols, figsize=(5,50))  
fig.tight_layout()  
  
for i in range(cols):  
    for j, row in data.iterrows():  
        x_selected = X_train[y_train == j]  
        axs[j][i].imshow(x_selected[random.randint(0,(len(x_selected) - 1)), :, :], cmap=plt.get_cmap('gray'))  
        axs[j][i].axis("off")  
        if i == 2:  
            axs[j][i].set_title(str(j) + " - " + row["SignName"])  
            num_of_samples.append(len(x_selected))  
print(num_of_samples)  
plt.figure(figsize=(12, 4))  
plt.bar(range(0, num_classes), num_of_samples)  
plt.title("Distribution of the train dataset")  
plt.xlabel("Class number")  
plt.ylabel("Number of images")  
plt.show()
```

0 - Speed limit (20km/h)



1 - Speed limit (30km/h)



2 - Speed limit (50km/h)



3 - Speed limit (60km/h)




4 - Speed limit (70km/h)



5 - Speed limit (80km/h)




6 - End of speed limit (80km/h)




7 - Speed limit (100km/h)



8 - Speed limit (120km/h)



9 - No passing



34 - Turn left ahead



35 - Ahead only



36 - Go straight or right



37 - Go straight or left



38 - Keep right




39 - Keep left



40 - Roundabout mandatory




41 - End of no passing



42 - End of no passing by vehicles over 3.5 metric tons



The bar chart below shows the number of images we have for each traffic sign category in our training dataset. Notice that the count is not uniform but let's just work with what we have.




Before we set up and train our model, we need to **preprocess the images** to ensure the best results. We want to **grayscale** them to simplify algorithms and reduce computational requirements.

```
[10] import cv2
```


```
plt.imshow(X_train[1000])  
plt.axis("off")  
print(X_train[1000].shape)  
print(y_train[1000])
```

```
(32, 32, 3)  
36
```



```
def grayscale(img):  
    img = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)  
    return img  
  
img = grayscale(X_train[1000])  
plt.imshow(img, cmap='gray_r')  
plt.axis("off")  
print(img.shape)
```

```
(32, 32)
```




We should also use the **cv2.equalizeHist()** method to spread out the pixel intensity values. This allows the image's areas with lower contrast to gain a higher contrast. See example below.

Next, we define a function called **preprocess()** that combines the **grayscale()** and **equalize()** methods, as well as normalise the pixel intensity value to a value between 0 and 1. This way, the numbers will be small and the computation becomes easier and faster.

```
#distributes the grayscale across each image uniformly
def equalize(img):
    img = cv2.equalizeHist(img)
    return img

img = equalize(img)
plt.imshow(img, cmap='gray_r')
plt.axis("off")
print(img.shape)
```

↳ (32, 32)



```
[13] #create function to apply image preprocessing to all images
def preprocess(img):
    img = grayscale(img)
    img = equalize(img)
    img = img/255
    return img

X_train = np.array(list(map(preprocess, X_train)))
X_val = np.array(list(map(preprocess, X_val)))
X_test = np.array(list(map(preprocess, X_test)))
```

We then apply the preprocess method on our feature datasets.



In order to input our feature data into the CNN model, we need to first **reshape our data** (which is current 32 pixel by 32 pixel) into this specific format (**32 pixel by 32 pixel by 1 channel**).

We also have to use the **to\_categorical()** method to **one-hot encode** our label datasets. One-hot encoding converts categorical information into a format that may be fed into machine learning algorithms to improve prediction accuracy.

```
▶ print(X_train.shape)
   print(X_test.shape)
   print(X_val.shape)
```

```
↳ (34799, 32, 32)
   (12630, 32, 32)
   (4410, 32, 32)
```

```
[16] #to input the data into CNN, need the depth dimension
      X_train = X_train.reshape(34799, 32, 32, 1)
      X_test = X_test.reshape(12630, 32, 32, 1)
      X_val = X_val.reshape(4410, 32, 32, 1)
```

```
[17] print(X_train.shape)

      (34799, 32, 32, 1)
```

```
[18] #one-hot encoding so that it can be fed into algorithm (increase prediction accuracy)
      y_train = to_categorical(y_train, 43)
      y_test = to_categorical(y_test, 43)
      y_val = to_categorical(y_val, 43)
```

Previously, the first image in our dataset has the label of 41 (referencing to End of No Passing road sign).

```
print(y_train[0])
```

```
41
```

After one-hot encoding, the label becomes a series of 43 binary values. Since the first image's label is essentially 41, notice that the 41<sup>st</sup> value is 1 while the rest are 0.

```
print(y_train[0])
```

```
[0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0.
 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 1. 0.]
```

A common trick in improving image recognition models is to **generate artificial images** to train the model. To do that, we can use the **Keras ImageDataGenerator** to translate the image vertically and horizontally, zoom into or out of the image, shear the image, and rotate the image.

We can set the batch size as 15 such that every time the generator is called, it will generate 15 new images.

```
[19] #generate additional artificial data of different variations (e.g. rotation, zoomed in/out, offset) to improve model performance
from keras.preprocessing.image import ImageDataGenerator

datagen = ImageDataGenerator(width_shift_range=0.1,          #max 10% shift
                             height_shift_range=0.1,        #max 10% shift
                             zoom_range=0.2,                #max 20% zoom in/out
                             shear_range=0.1,               #max 10% shear
                             rotation_range=10)             #max 10degrees rotation

datagen.fit(X_train)


#to get datagen to start generating new images
#batches is an iterator object, can be called by next() method (will generate 15 images everytime method is called)
batches = datagen.flow(X_train, y_train, batch_size = 15)
X_batch, y_batch = next(batches)

fig, axs = plt.subplots(1, 15, figsize=(20, 5))
fig.tight_layout()

for i in range(15):
    axs[i].imshow(X_batch[i].reshape(32, 32))
    axs[i].axis("off")

print(X_batch.shape)
```

(15, 32, 32, 1)



Next, let's create the **CNN model**. There is no specific formula for creating a model and is usually done through a series of trial and error. In Convolutional Neural Networks, **filters** (e.g. 60, 30 below) detect spatial patterns such as edges in an image by detecting the changes in intensity values of the image.

**Kernel size** (e.g. (5, 5) or (3, 3) below) is the size of the filter.

```
#create model
def modified_model():
    model = Sequential()
    model.add(Conv2D(60, (5, 5), input_shape=(32, 32, 1), activation='relu'))
    model.add(Conv2D(60, (5, 5), activation='relu'))
    model.add(MaxPooling2D(pool_size=(2, 2)))

    model.add(Conv2D(30, (3, 3), activation='relu'))
    model.add(Conv2D(30, (3, 3), activation='relu'))
    model.add(MaxPooling2D(pool_size=(2, 2)))

    model.add(Flatten())
    model.add(Dense(500, activation='relu'))
    model.add(Dropout(0.5))
    model.add(Dense(43, activation='softmax'))

    model.compile(Adam(lr = 0.001), loss='categorical_crossentropy', metrics=['accuracy'])
    return model

model = modified_model()
print(model.summary())
```

For a multi-classification problem, we can set our last **activation function as 'softmax'**. The softmax activation function transforms the raw outputs of the neural network into a vector of probabilities, allowing us to return the classification with the highest probability.

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 28, 28, 60)	1560
conv2d_1 (Conv2D)	(None, 24, 24, 60)	90060
max_pooling2d (MaxPooling2D)	(None, 12, 12, 60)	0
conv2d_2 (Conv2D)	(None, 10, 10, 30)	16230
conv2d_3 (Conv2D)	(None, 8, 8, 30)	8130
max_pooling2d_1 (MaxPooling2D)	(None, 4, 4, 30)	0
flatten (Flatten)	(None, 480)	0
dense (Dense)	(None, 500)	240500
dropout (Dropout)	(None, 500)	0
dense_1 (Dense)	(None, 43)	21543
=====		
Total params: 378,023		
Trainable params: 378,023		
Non-trainable params: 0		

Now that we have set up our model, let's **fit the model** with our training images. Since we have an image generator, we can use the `model.fit_generator()` method to create new images and to fit the model.


**Batch size** refers to the number of new images we generate each time we call the generator and **steps\_per\_epoch** refers to the number of times we call the generator in 1 epoch. An **epoch** just basically means when all training datasets have been passed through the model once.

```
history = model.fit_generator(datagen.flow(X_train, y_train, batch_size=50),
                             steps_per_epoch=X_train.shape[0]/50,
                             epochs=10,
                             validation_data=(X_val, y_val), shuffle = 1)
```

```
Epoch 1/10
<ipython-input-22-7dbfefff6b2a>:1: UserWarning: `Model.fit_generator` is deprecated and will be removed in a future version. Please u
  history = model.fit_generator(datagen.flow(X_train, y_train, batch_size=50),
695/695 [=====] - 21s 22ms/step - loss: 1.7361 - accuracy: 0.5076 - val_loss: 0.3688 - val_accuracy: 0.8866
Epoch 2/10
695/695 [=====] - 15s 21ms/step - loss: 0.5454 - accuracy: 0.8339 - val_loss: 0.1271 - val_accuracy: 0.9605
Epoch 3/10
695/695 [=====] - 15s 21ms/step - loss: 0.3171 - accuracy: 0.9015 - val_loss: 0.1141 - val_accuracy: 0.9658
Epoch 4/10
695/695 [=====] - 15s 22ms/step - loss: 0.2404 - accuracy: 0.9248 - val_loss: 0.0547 - val_accuracy: 0.9830
Epoch 5/10
695/695 [=====] - 15s 22ms/step - loss: 0.1985 - accuracy: 0.9369 - val_loss: 0.0469 - val_accuracy: 0.9859
Epoch 6/10
695/695 [=====] - 16s 24ms/step - loss: 0.1762 - accuracy: 0.9466 - val_loss: 0.0606 - val_accuracy: 0.9796
Epoch 7/10
695/695 [=====] - 15s 22ms/step - loss: 0.1546 - accuracy: 0.9523 - val_loss: 0.0536 - val_accuracy: 0.9837
Epoch 8/10
695/695 [=====] - 17s 24ms/step - loss: 0.1348 - accuracy: 0.9586 - val_loss: 0.0671 - val_accuracy: 0.9766
Epoch 9/10
695/695 [=====] - 15s 22ms/step - loss: 0.1255 - accuracy: 0.9616 - val_loss: 0.0443 - val_accuracy: 0.9846
Epoch 10/10
695/695 [=====] - 15s 22ms/step - loss: 0.1172 - accuracy: 0.9638 - val_loss: 0.0346 - val_accuracy: 0.9893
```

```
plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.legend(['training', 'test'])
plt.title('Loss')
plt.xlabel('epoch')
```


Text(0.5, 0, 'epoch')



Notice that the validation accuracy is higher than the training accuracy. Technically, we can continue to run the model for a few more epochs to see if they converge.

```
plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.legend(['training', 'test'])
plt.title('Accuracy')
plt.xlabel('epoch')
```

Text(0.5, 0, 'epoch')



Next, we measure how our model work on the test data (this is the 3<sup>rd</sup> set of data that the model had not seen before). Great, **96.7% accuracy!**

```
[25] #Evaluate model on test data
      score = model.evaluate(X_test, y_test, verbose=0)
      print('Test score:', score[0])
      print('Test accuracy:', score[1])
```

```
Test score: 0.1292283833026886
Test accuracy: 0.9671417474746704
```

Now that we have a trained model, let's **test it out** on random images culled from the internet. Yay, the model works well!

```
#predict never seen before images on internet
import requests
from PIL import Image
url = 'https://c8.alamy.com/comp/A0RX23/cars-a'
r = requests.get(url, stream=True)
img = Image.open(r.raw)
plt.imshow(img, cmap=plt.get_cmap('gray'))
```


<matplotlib.image.AxesImage at 0x7f8613b4b130>



```
img = np.asarray(img)
img = cv2.resize(img, (32, 32))
img = preprocess(img)
plt.imshow(img, cmap = plt.get_cmap('gray'))
print(img.shape)
img = img.reshape(1, 32, 32, 1)


print("predicted sign: " + str(np.argmax(model.predict(img), axis = 1)))
print("predicted sign: " + data['SignName'].iloc[np.argmax(model.predict(img), axis = 1)])
```

```
(32, 32)
1/1 [=====] - 0s 150ms/step
predicted sign: [34]
1/1 [=====] - 0s 18ms/step
34 predicted sign: Turn left ahead
Name: SignName, dtype: object
```






1 predicted sign: Speed limit (30km/h)  
Name: SignName, dtype: object







23    predicted sign: Slippery road  
Name: SignName, dtype: object








13    predicted sign: Yield  
Name: SignName, dtype: object





29 predicted sign: Bicycles crossing  
Name: SignName, dtype: object





14    predicted sign: Stop  
Name: SignName, dtype: object

